

Fault tolerance becomes critical at Petascale (MTTI  $\leq 1$  day)  
Poor fault tolerance design may lead to huge overhead

Overhead of checkpoint/restart

Cost of non optimal checkpoint intervals:

100%  
0%

Today, 20% or more of the computing capacity in a large high-performance computing system is wasted due to failures and recoveries.

Dr. E.N. (Mootaz) Elnozahy et al. *System Resilience at Extreme Scale*,  
DARPA

